

# CSC / CPE 357

Systems Programming

# Topics

---

- File I/O

# Standard I/O: File Access

---

- Examples thus far have relied on `stdin/stdout` which are made available automatically
- To read or write to a named file using standard I/O functions, a file must first be opened by the library function `fopen()`
  - `FILE *fopen(char *name, char *mode)`
- `FILE *` (file pointer) represents a structure that contains information about the file, allowing other functions to read/write
- `mode` a character string indicating the intended use of the file:
  - `"r"` read
  - `"w"` write
  - `"a"` append
  - (additional options described here: `man 3 fopen`)

# Standard I/O: File Access Modes

---

Read Modes	Write Modes	Append Modes
<ul style="list-style-type: none"><li>• <code>r</code>: open for reading</li><li>• <code>rb</code>: open for reading in binary mode</li><li>• <code>r+</code>: open for reading or writing</li><li>• <code>rb+</code>: binary mode read or write</li></ul> <p><code>fopen()</code> returns <code>NULL</code> if file does not exist</p>	<ul style="list-style-type: none"><li>• <code>w</code>: write</li><li>• <code>wb</code>: write, binary mode</li><li>• <code>w+</code>: read or write</li><li>• <code>wb+</code>: read or write, binary mode</li></ul> <p>If the file exists, contents are overwritten. File is created if it does not exist.</p>	<ul style="list-style-type: none"><li>• <code>a</code>: open for appending</li><li>• <code>ab</code>: append, binary mode</li><li>• <code>a+</code>: append or read</li><li>• <code>ab+</code>: append or read binary mode</li></ul> <p>New data will be added to the end of the file. File is created if it does not exist.</p>

# Standard I/O: File Access

---

- `fopen()` will fail (returning `NULL`) if
  - In read mode, the file does not exist
  - The current user does not have permission to access the file
- If `fopen()` succeeds, characters can be read/written using:
  - `int getc(FILE *fp)`
  - `int putc(int c, FILE *fp)`
- Formatted read/write based on a file pointer:
  - `int fprintf(FILE *fp, char *format, ...)`
  - `int fscanf(FILE *fp, char *format, ...)`
- When finished reading/writing:
  - `int fclose(FILE *fp)`

# Example Code

---

In class example: Copying contents from one file to another

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main() {
5      // Open the source file in read mode
6      FILE *sourceFile = fopen("input.txt", "r");
7      if (sourceFile == NULL) {
8          perror("Error opening source file");
9          return 1;
10     }
11
12     // Open the destination file in write mode
13     FILE *destinationFile = fopen("output.txt", "w");
14     if (destinationFile == NULL) {
15         perror("Error opening destination file");
16         // Close the source file before exiting
17         fclose(sourceFile);
18         return 1;
19     }
20
```

```
21     int ch; // Variable to store each character
22
23     // Read from source file character by character using getc
24     while ((ch = getc(sourceFile)) != EOF) {
25         // Write the character to the destination file using putc
26         putc(ch, destinationFile);
27     }
28
29     // Close both files
30     fclose(sourceFile);
31     fclose(destinationFile);
32
33     printf("File copied.\n");
34
35     return 0;
36 }
```

## Use of *getline()*

---

- The `getline()` function reads a line from stream, delimited by the character newline.
- The `getline()` functions return the number of characters written, excluding the terminating NUL character. The value -1 is returned if an error occurs, or if end-of-file is reached.

```
ssize_t getline(char ** restrict linep, size_t * restrict linecapp, FILE *  
restrict stream);
```

# Example Code

## In class example: Lab 2 task 6

```
1  #define _GNU_SOURCE
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <string.h>
5
6  int main(int argc, char *argv[]) {
7      // Ensure a file name is provided as an argument
8      if (argc != 2) {
9          fprintf(stderr, "Usage: %s <filename>\n", argv[0]);
10         return 1;
11     }
12
13     // Open the file
14     FILE *file = fopen(argv[1], "r");
15     if (file == NULL) {
16         perror("Error opening file");
17         return 1;
18     }
19
20     char *line = NULL;
21     unsigned int len = 0;
22     int read;
23
24     // Variables to store the last two lines
25     char *last_line1 = NULL;
26     char *last_line2 = NULL;
27
```

```
28     // Read the file line by line
29     while ((read = getline(&line, (size_t *)&len, file)) != -1) {
30         // Free the second last line to avoid memory leaks
31         if (last_line2) {
32             free(last_line2);
33         }
34
35         // Move last_line1 to last_line2 and
36         // store the current line in last_line1
37         last_line2 = last_line1;
38         last_line1 = strdup(line); // Make a copy of the line
39     }
40
41     // Print the last two lines if they exist
42     if (last_line2) {
43         printf("%s", last_line2);
44     }
45     if (last_line1) {
46         printf("%s", last_line1);
47     }
48
49     // Free the memory and close the file
50     free(last_line1);
51     free(last_line2);
52     free(line);
53     fclose(file);
54
55     return 0;
56 }
```



# Standard I/O: Line Input and Output

---

- To read an entire line from a file (including the newline):
  - `char *fgets(char *line, int maxline, FILE *fp)`
  - Reads the next input line into the character array `line`
  - Returns the line read or NULL on end of file or error
- Write a string to a file (with or without a newline):
  - `int fputs(char *line, FILE *fp)`

# Unbuffered I/O

---

Open via system calls:

```
int open(const char *path, int oflag, ...);
```

Arguments:

- `O_RDONLY`      open for reading only
- `O_WRONLY`      open for writing only
- `O_RDWR`      open for reading and writing
- `O_SEARCH`      open directory for searching
- `O_EXEC`      open for execute only

Documentation: `man 2 open`

(recall: manual page section 2 contains information about system calls)

# Unbuffered I/O

---

Input and output via system calls:

```
int n_read = read(int fd, char *buf, int n);  
int n_written = write(int fd, char *buf, int n);
```

Arguments:

- fd: file descriptor
- buf: buffer where data is read/written
- n: number of bytes to be read/written

Documentation: `man 2 read`

(recall: manual page section 2 contains information about system calls)

# File Descriptors

---

- A **file descriptor** is a non-negative integer.
  - Kernel returns a file descriptor on `open()` of an existing file or `creat()` a new file
- UNIX Conventions
  - File descriptor 0 represents `STDIN`
  - File descriptor 1 represents `STDOUT`
  - File descriptor 2 represents `STDERR`
- `STDIN_FILENO`, `STDOUT_FILENO`, and `STDERR_FILENO` constants are defined in `<unistd.h>`
- Contrast with *file pointers* `stdin`, `stdout`, `stderr`, declared in `<stdio.h>`
  - `int fileno(FILE *stream)`
  - `FILE *fdopen(int fd, const char *mode);`