

CSC / CPE 357

Systems Programming

Integer Data Types in C

```
#include <stdint.h>
```

	Signed	Unsigned	Maximum Value (Unsigned)
?-Bit Integer (at least 16 bits according to C specification)	int	unsigned int	?
8-bit Integer	int8_t	uint8_t	UINT8_MAX ($2^8 - 1$)
16-bit Integer	int16_t	uint16_t	UINT16_MAX
32-bit Integer	int32_t	uint32_t	UINT32_MAX
64-bit Integer	int64_t	uint64_t	UINT32_MAX

Number Representation in C

- In C, values can be specified in octal or hexadecimal instead of decimal.
- Leading 0 (zero) on an integer constant means octal
 - `int i = 037; // octal 37, decimal 31`
- Leading 0x or 0X means hexadecimal
 - `int i = 0x1f; // hexadecimal 1f, decimal 31`
- Octal and hexadecimal constants may also be followed by:
 - L for long
 - U for unsigned
 - 0XFUL is an unsigned long constant with value 15 decimal.

[Handy Conversion Tool](#)

Bitwise Operators in C

C provides six operators for bit manipulation; these may only be applied to integral operands: char, short, int, and long (signed or unsigned)

Symbol	Operation
&	bitwise AND
	bitwise OR
^	bitwise XOR
~	bitwise NOT
<<	left shift
>>	right shift

Bitwise vs Logical Operators

- Easy to confuse two different sets of operators in C:
 - Bitwise operators: `&` and `|`
 - Logical operators: `&&` and `||` (logical AND, logical OR)
 - Left-to-right evaluation of a truth value
- If `x` is 1 and `y` is 2:
 - `x & y` is zero
 - `x && y` is one

Bitwise AND

The bitwise AND operator & is often used to mask bits

```
n = n & 0177; // octal 177, dec 127, hex 0x7f, bin 01111111
```

sets to zero all but the low-order 7 bits of n.

		A & B
0	1	0
0	0	0
1	1	1
1	0	0

Bitwise OR

The bitwise OR operator `|` is used to turn bits on:

```
x = x | SET_ON;
```

sets to one in `x` the bits that are set to one in `SET_ON`.

		A B
0	1	1
0	0	0
1	1	1
1	0	1

Bitwise XOR

The bitwise exclusive OR operator \wedge sets a one in each bit position where its operands have different bits, and zero where they are the same.

		$A \wedge B$
0	1	1
0	0	0
1	1	0
1	0	1

Bitwise NOT

The unary operator `~` yields the one's complement of an integer -- converts each 1-bit into a 0-bit and vice versa.

For example:

```
x = x & ~077
```


sets the last six bits of x to zero

Shift Operators


- The shift operators `<<` and `>>` perform left and right shifts
 - Number of bit positions given by the right operand, must be non-negative.
- For example, `x << 2` shifts the value of `x` to the left by two bits
 - Vacated bits filled with zero
 - `x << 2` is equivalent to multiplication by 4
- Right shifting *unsigned* value always fits the vacated bits with zero.
- Right shifting *signed* value is implementation-dependent:
 - Fill with bit signs ("arithmetic shift") on some machines/compiler (x86/gcc)
 - Fill with 0 ("logical shift") on others

Shift Operators: Arithmetic vs Logical


- **Arithmetic shift:** sign bit copied to preserve sign
- **Logical shift:** zeros are shifted in to replace the discarded bits



Arithmetic Right Shift



Logical Right Shift



Left Shift (Logical and Arithmetic are Identical)

Negative Numbers

- In the C programming language, signed integers are represented using **Two's Complement**.
- Most significant bit of the number (also called MSB) is used as the sign
 - 0 for positive numbers
 - 1 for negative numbers.
- Remaining bits are used to store the number's value.
- 0 is considered a positive number

Two's Complement

- We create a Two's Complement representation of a negative number by inverting all bits (recall the `~` operator) of the positive value, then adding 1.
- For example, binary representation of -47 (decimal):
 - Binary representation of *positive* 47: 00101111
 - Flip all bits: 11010000
 - Add 1: 11010001
- This representation limits our range of values
 - Maximum positive value signed `int8_t` can hold is: $2^7 - 1 = 127$
 - Minimum value: -2^7 (since 0 is considered positive)

Bit Masking

Bit masking refers to operations like the following:

- Set a specific bit to a value (0 or 1)
- Clear a specific bit, or a sequence of bits
- Invert the values of all bits

Example: File Permissions

Recall chmod (change file permissions):

```
#include <sys/stat.h>
```

```
int chmod(const char *pathname, mode_t mode);
```

The mode argument is a bitmask. To change permissions, there are several defined constants that can be applied along with bitwise operators.

Constant	Description	Effect on regular file	Effect on directory
S_ISUID	set-user-ID	set effective user ID on execution	(not used)
S_ISGID	set-group-ID	if group-execute set, then set effective group ID on execution; otherwise, enable mandatory record locking (if supported)	set group ID of new files created in directory to group ID of directory
S_ISVTX	sticky bit	control caching of file contents (if supported)	restrict removal and renaming of files in directory
S_IRUSR	user-read	user permission to read file	user permission to read directory entries
S_IWUSR	user-write	user permission to write file	user permission to remove and create files in directory
S_IXUSR	user-execute	user permission to execute file	user permission to search for given pathname in directory
S_IRGRP	group-read	group permission to read file	group permission to read directory entries
S_IWGRP	group-write	group permission to write file	group permission to remove and create files in directory
S_IXGRP	group-execute	group permission to execute file	group permission to search for given pathname in directory
S_IROTH	other-read	other permission to read file	other permission to read directory entries
S_IWOTH	other-write	other permission to write file	other permission to remove and create files in directory
S_IXOTH	other-execute	other permission to execute file	other permission to search for given pathname in directory

Figure 4.26 Summary of file access permission bits

Example: The Lonely Integer

Given an array of integer values, where all elements except one occur twice, find the unique element, the so-called lonely integer.

For example, if $L = \{1, 2, 3, 3, 8, 1, 9, 2, 9\}$ the unique element is 8, because the rest of the elements come in pairs.

Example: `getbits()`

`getbits(x, p, n)` returns the n -bit field of x that begins at position p .

Assume:

- bit position 0 is at the right end
- n and p are positive values, within range

For example, `getbits(x, 4, 3)` returns the three bits in positions 4, 3 and 2 from the right