

# CSC / CPE 357

## Systems Programming

Chapter 17 in Advanced Programming  
in the UNIX Environment

# IPC Summary

---

- UNIX supports several mechanisms for interprocess communication:
  - Pipes
  - Named pipes (FIFOs)
  - Other forms IPC commonly called "XSI IPC"
    - message queues
    - semaphores (a synchronization primitive)
    - shared memory
- Message queues are rarely used due to their complexity and lack of any performance advantage
- **Sockets** provide another means of communication between processes (either on the same machine or separate machines)

# UNIX Domain Sockets


---

- The capabilities of **UNIX Domain Sockets** fall between pipes and network sockets
  - Full duplex pipe using `socketpair()`
  - Local-only socket using sockets API

# socketpair()

---

The `socketpair()` function creates a pair of connected UNIX domain sockets that represent a full-duplex pipe: both ends are open for reading and writing



**Figure 17.1** A socket pair

```
int socketpair(int domain, int type, int protocol, int sv[2]);
```

Returns 0 if OK, -1 on error

## socketpair() example

---

socketpair.c

Note that socketpair can only be used between related processes (similar to `pipe()`, but full duplex)

# Named UNIX Domain Sockets

---

UNIX domain sockets can be either unnamed (ie. `socketpair()`) or bound to a filesystem pathname.

UNIX domain socket address is represented by the structure:

```
struct sockaddr_un {
    sa_family_t sun_family;           /* AF_UNIX */
    char        sun_path[108];        /*
Pathname */
};
```

# Socket Types

---

`SOCK_DGRAM` - message/datagram socket that preserves message boundaries (with UNIX domain sockets, this is reliable and does not reorder datagrams)

`SOCK_STREAM` - stream socket, message boundaries not preserved

`SOCK_SEQPACKET` - sequenced-packet socket that is connection-oriented, preserves message boundaries, and delivers messages in the order that they were sent.

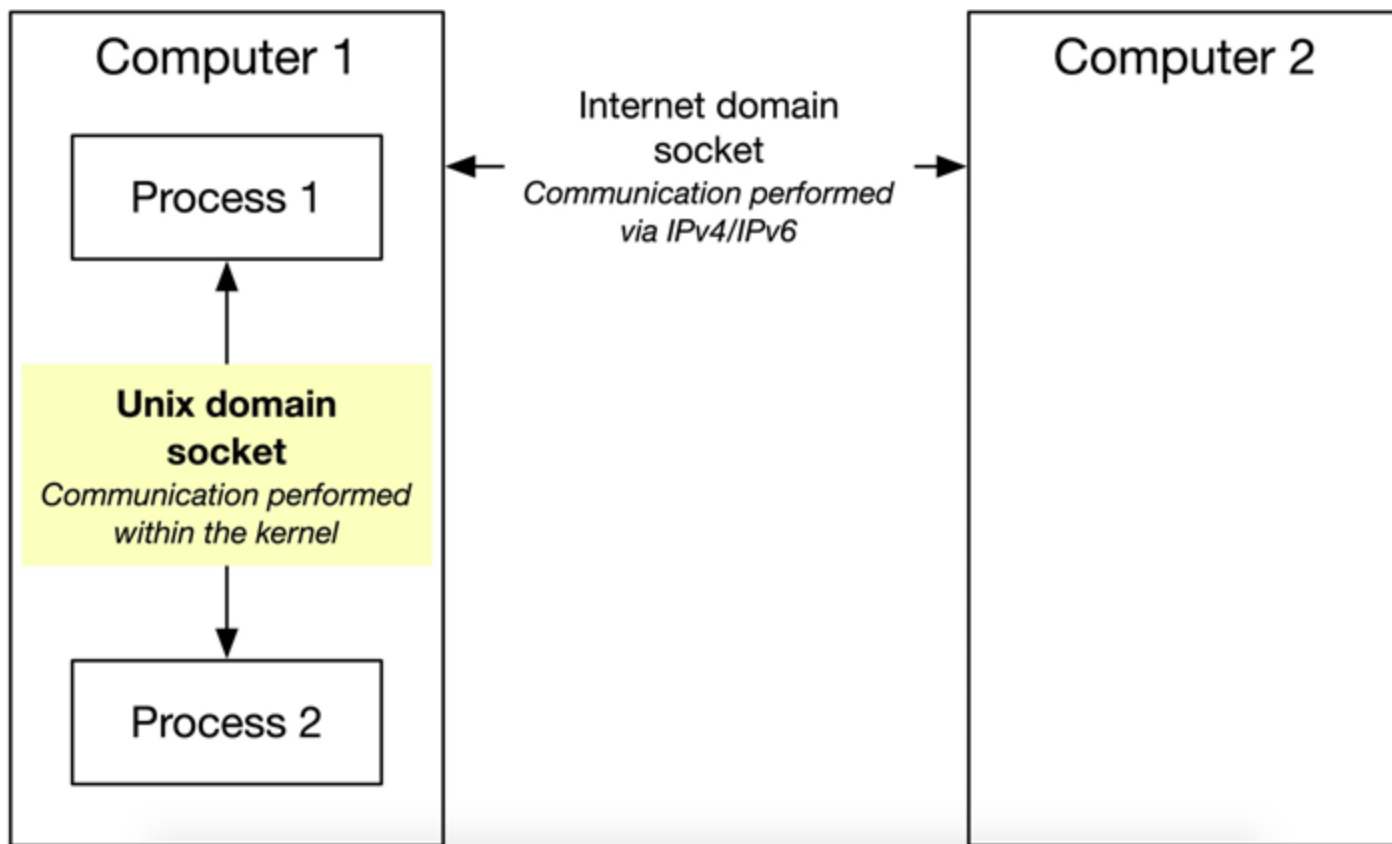
# UNIX Domain Sockets Example

---

`unix_server.c`

`unix_client.c`





# UNIX Domain Sockets

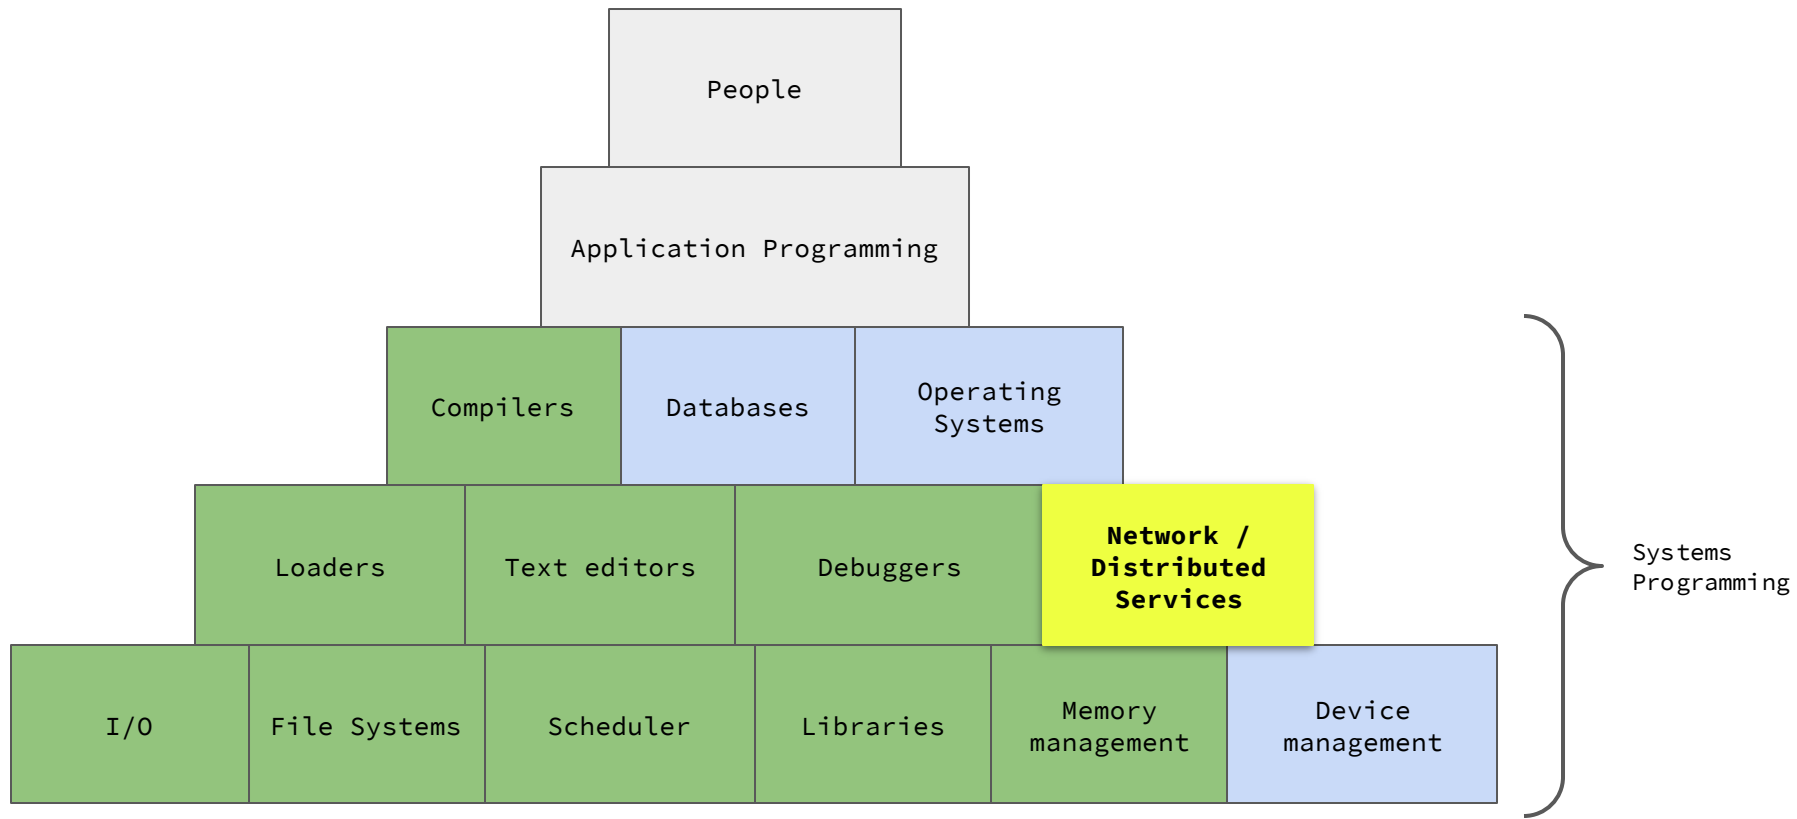
---

- UNIX domain sockets may be used to communicate between processes running on the same machine.
- Although Internet sockets can be used for this same purpose, UNIX domain sockets are more efficient:
  - no protocol processing to perform,
  - no network headers to add or remove,
  - no checksums to calculate,
  - no sequence numbers to generate,
  - no acknowledgements to send.

# IPC Summary

---

- pipe()
  - half duplex (i.e., one way communication) between related processes
- FIFO (aka Named Pipe)
  - half-duplex, exists as a file, can be used between unrelated processes
- UNIX Domain Sockets
  - Local-only socket using sockets API
  - Less overhead (higher performance) when compared to Internet Sockets
- Internet Sockets
  - TCP
    - i. connects any two processes including remote processes
    - ii. high protocol overhead
    - iii. reliable byte stream, but message boundaries are not preserved
  - UDP
    - i. lower protocol overhead than TCP
    - ii. message boundaries are preserved, but deliveries are unreliable



# Network Communication

---

- Socket IPC interface can be used to communicate between processes on different machines
- Communications are governed by protocols
  - TCP, UDP, IP
  - Higher-level protocols (email/SMTP, web/HTTP) implemented on top of fundamental protocols

# Sockets

---

- A **socket** is an abstraction of a communication endpoint
- Socket descriptors are implemented as file descriptors in the UNIX System
- Many functions that deal with file descriptors will work with a socket descriptor
  - For example: `read()`, `write()`

# socket() function

---

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol);
```

Returns: file (socket) descriptor if OK, -1 on error

- **domain** determines the nature of the communication, including the address format
  - AF\_INET - IPv4 Internet domain
  - AF\_INET6 - IPv6 Internet domain
- **type** specifies communication details
  - SOCK\_DGRAM - fixed-length, connectionless, unreliable messages
  - SOCK\_STREAM - sequenced, reliable, bidirectional, connection-oriented byte streams

# shutdown()

---

```
#include <sys/socket.h>
```

```
int shutdown(int sockfd, int how);
```

Returns: 0 if OK, -1 on error

- **how** may be one of: SHUT\_RD, SHUT\_WR or SHUT\_RDWR
- Why not simply close() as we do with file descriptors?
  - If we dup() the socket it won't be deallocated until the last file descriptor referring to it is closed; shutdown deactivates socket entirely
  - Sometimes convenient to shut a socket down in one direction only



## Addresses / Ports

---

- Internet protocols use an IP **address** and a **port** to identify an endpoint
- To connect to another machine you need to be able to find it (by address)
- To connect to a particular program on another machine you need to be able to differentiate it (by port)

## Address Formats

---

In the IPv4 Internet domain (AF\_INET), a socket address is represented by the `sockaddr_in` structure (which embeds `in_addr`):

```
struct in_addr {  
    in_addr_t      s_addr;          /* IPv4 address */  
};  
  
struct sockaddr_in {  
    sa_family_t     sin_family;     /* address family */  
    in_port_t       sin_port;       /* port number */  
    struct in_addr   sin_addr;      /* IPv4 address */  
};
```

# Human-Readable Addresses

---

The `inet_ntop` function converts a binary address in network byte order into a human-readable text string; `inet_pton` converts a text string into a binary address in network byte order.

```
#include <arpa/inet.h>
```

```
const char *inet_ntop(int domain, const void *restrict addr,  
                      char *restrict str, socklen_t size);
```

Returns: pointer to address string on success, NULL on error

```
int inet_pton(int domain, const char *restrict str,  
             void *restrict addr);
```

Returns: 1 on success, 0 if the format is invalid, or -1 on error

# Associating Addresses with Sockets

---

For a server, we need to associate a well-known address with the server's socket on which client requests will arrive.

The `bind` function associates an address with a socket

```
#include <sys/socket.h>
```

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t len);
```

Returns: 0 if OK, -1 on error

The address we specify must be valid for the machine and the address family we used to create the socket

# Host / Network Byte Order

---

- Network protocols specify a byte ordering to allow different computer systems can exchange protocol information without confusion

- TCP/IP protocol suite uses big-endian byte order

```
#include <arpa/inet.h>
```

```
uint32_t htonl(uint32_t hostint32);
```

Returns: 32-bit integer in network byte order

```
uint16_t htons(uint16_t hostint16);
```

Returns: 16-bit integer in network byte order

```
uint32_t ntohl(uint32_t netint32);
```

Returns: 32-bit integer in host byte order

```
uint16_t ntohs(uint16_t netint16);
```

Returns: 16-bit integer in host byte order

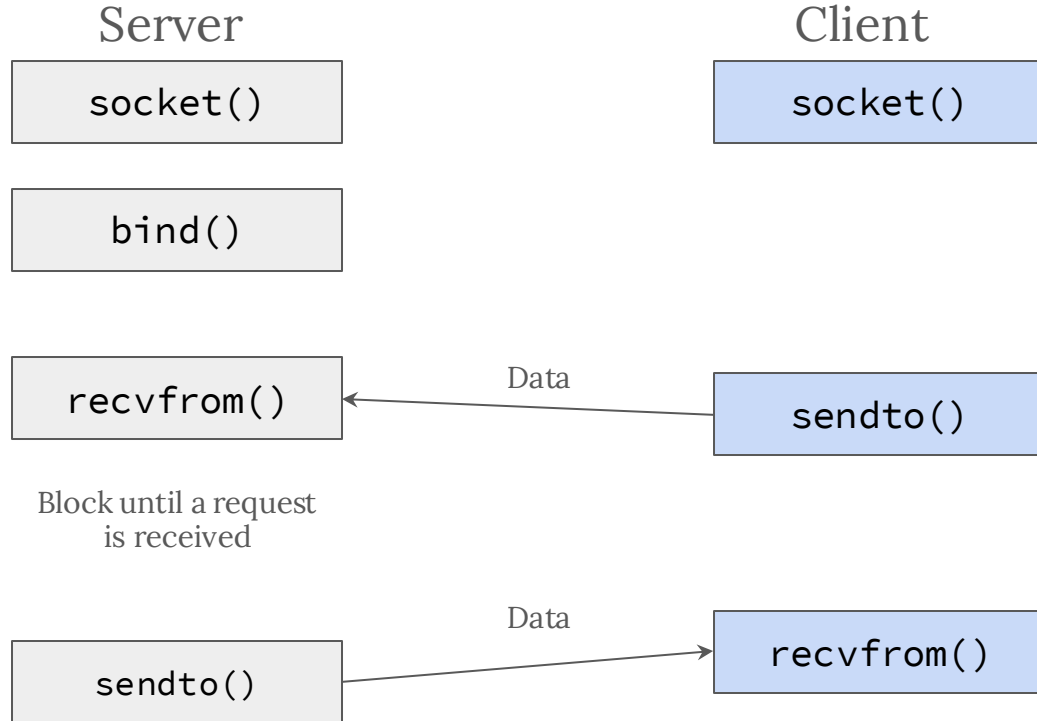
# User Datagram Protocol (UDP)

---

- UDP is connectionless and unreliable
  - Create a socket, send a message
  - No guarantee of successful or complete delivery
- Similar to dropping a note in regular postal mail

# Datagram (UDP) Socket Communication

---



# UDP Example

---

`udp_server.c`

`udp_client.c`



# Transmission Control Protocol (TCP)

---

- TCP is connection-oriented and reliable
  - Server must listen and accept client requests
  - Client and server must establish a connection before data can be sent
  - Error detection (at the expense of higher latency)
- Similar to sending a letter with delivery confirmation

# listen() for requests

---

A server announces that it is willing to accept connect requests by calling the listen function.

```
#include <sys/socket.h>

int listen(int sockfd, int backlog);
```

Returns: 0 if OK, -1 on error

backlog argument provides a hint to the system regarding the number of outstanding connect requests that it should enqueue on behalf of the process

## accept() a request

---

The `accept` function retrieves a connect request and converts it into a connection

```
#include <sys/socket.h>
```

```
int accept(int sockfd, struct sockaddr *restrict addr,  
           socklen_t *restrict len);
```

Returns: file (socket) descriptor if OK, -1 on error

File descriptor returned by `accept` is a socket descriptor that is connected to the client that called `connect`

## accept() – Retrieve Client Identity

---

```
#include <sys/socket.h>
```

```
int accept(int sockfd, struct sockaddr *restrict addr,  
           socklen_t *restrict len);
```

Returns: file (socket) descriptor if OK, -1 on error

If we don't care about the client's identity, we can set the `addr` and `len` parameters to `NULL`

If `sockaddr` / `socklen_t` pointers are provided, `accept` will fill in the client's address size of the address.

# Data Transfer

---

- Possible to use `read()` and `write()` with sockets
  - Limited to simple data transfer
- Specialized socket functions allow:
  - Specification of options
  - Receipt of packets from multiple clients

```
#include <sys/socket.h>
```

```
ssize_t send(int sockfd, const void *buf, size_t nbytes, int flags);
```

Returns: number of bytes sent if OK, -1 on error

```
#include <sys/socket.h>
```

```
ssize_t recv(int sockfd, void *buf, size_t nbytes, int flags);
```

Returns: length of message in bytes,  
0 if no messages are available and peer has done an orderly shutdown,  
or -1 on error

# TCP Example

---

`tcp_server.c`

`tcp_client.c`

# Stream (TCP) Socket Communication

---

